

Bayesian Learning

Based on “Machine Learning”, T. Mitchell, McGRAW Hill, 1997, ch. 6

Acknowledgement:

The present slides are an adaptation of slides drawn by T. Mitchell

Two Roles for the Bayesian Methods in Learning

1. Provides practical **learning algorithms**

by combining prior knowledge/probabilities with observed data:

- Naive Bayes learning algorithm
- Expectation Maximization (EM) learning algorithm (scheme): learning in the presence of unobserved variables
- Bayesian Belief Network learning

2. Provides a **useful conceptual framework**

- Serves for evaluating other learning algorithms, e.g. concept learning through general-to-specific hypotheses ordering (FINDS, and CANDIDATEELIMINATION), neural networks, linear regression
- Provides additional insight into Occam's razor

PLAN

2.

1. Basic Notions

Bayes' Theorem

Defining **classes of hypotheses**:

Maximum A posteriori Probability (MAP) hypotheses

Maximum Likelihood (ML) hypotheses

2. Learning MAP hypotheses

2.1 The brute force MAP hypotheses learning algorithm

2.2 The Bayes optimal classifier;

2.3 The Gibbs classifier;

2.4 The **Naive Bayes** and the **Joint Bayes** classifiers.

Example: Learning over text data using Naive Bayes

2.5 The Minimum Description Length (MDL) Principle;
MDL hypotheses

3. Learning ML hypotheses

3.1 ML hypotheses in learning real-valued functions

3.2 ML hypotheses in learning to predict probabilities

3.3 **The Expectation Maximization (EM) algorithm**

4. **Bayesian Belief Networks**

1 Basic Notions

- **Product Rule:**

probability of a conjunction of two events A and B:

$$P(A \wedge B) = P(A|B)P(B) = P(B|A)P(A)$$

- **Bayes' Theorem:**

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

- **Theorem of total probability:**

if events A_1, \dots, A_n are mutually exclusive,
with $\sum_{i=1}^n P(A_i) = 1$, then

$$P(B) = \sum_{i=1}^n P(B|A_i)P(A_i)$$

in particular

$$P(B) = P(B|A)P(A) + P(B|\neg A)P(\neg A)$$

Using Bayes' Theorem for Hypothesis Learning

$$P(h|D) = \frac{P(D|h)P(h)}{P(D)}$$

- $P(D)$ = the (prior) probability of training data D
- $P(h)$ = the (prior) probability of the hypothesis h
- $P(D|h)$ = the (a posteriori) probability of D given h
- $P(h|D)$ = the (a posteriori) probability of h given D

Classes of Hypotheses

Maximum Likelihood (ML) hypothesis:

the hypothesis that best explains the training data

$$h_{ML} = \operatorname{argmax}_{h_i \in H} P(D|h_i)$$

Maximum A posteriori Probability (MAP) hypothesis:

the most probable hypothesis given the training data

$$h_{MAP} = \operatorname{argmax}_{h \in H} P(h|D) = \operatorname{argmax}_{h \in H} \frac{P(D|h)P(h)}{P(D)} = \operatorname{argmax}_{h \in H} P(D|h)P(h)$$

Note: If $P(h_i) = P(h_j), \forall i, j$, then $h_{MAP} = h_{ML}$

Exemplifying MAP Hypotheses

Suppose the following data characterize the lab result for cancer-suspect people.

$$\begin{array}{cc|c} P(\text{cancer}) = 0.008 & P(\neg\text{cancer}) = 0.992 & h_1 = \text{cancer}, h_2 = \neg\text{cancer} \\ \hline P(+|\text{cancer}) = 0.98 & P(-|\text{cancer}) = 0.02 & D = \{+, -\}, P(D | h_1), P(D | h_2) \\ P(+|\neg\text{cancer}) = 0.03 & P(-|\neg\text{cancer}) = 0.97 & \end{array}$$

Question: Should we diagnose a patient x whose lab result is positive as having cancer?

Answer: No.

Indeed, we have to find $\text{argmax}\{P(\text{cancer}|+), P(\neg\text{cancer}|+)\}$.

Applying Bayes theorem (for $D = \{+\}$):

$$\left. \begin{array}{l} P(+ | \text{cancer})P(\text{cancer}) = 0.98 \times 0.008 = 0.0078 \\ P(+ | \neg\text{cancer})P(\neg\text{cancer}) = 0.03 \times 0.992 = 0.0298 \end{array} \right\} \Rightarrow h_{MAP} = \neg\text{cancer}$$

(We can infer $P(\text{cancer} | +) = \frac{0.0078}{0.0078+0.0298} = 21\%$)

2 Learning MAP Hypothesis

2.1 The Brute Force MAP Hypothesis Learning Algorithm

Training:

Choose the hypothesis with the highest posterior probability

$$h_{MAP} = \operatorname{argmax}_{h \in H} P(h|D) = \operatorname{argmax}_{h \in H} P(D|h)P(h)$$

Testing:

Given x , compute $h_{MAP}(x)$

Drawback:

Requires to compute all probabilities $P(D|h)$ and $P(h)$.

2.2 The Bayes Optimal Classifier:

The Most Probable Classification of New Instances

So far we've sought h_{MAP} , the *most probable hypothesis* given the data D .

Question: Given new instance x — the classification of which can take any value v_j in some set V —, what is its *most probable classification*?

Answer: $P(v_j|D) = \sum_{h_i \in H} P(v_j|h_i)P(h_i|D)$

Therefore, the **Bayes optimal classification** of x is:

$$\operatorname{argmax}_{v_j \in V} \sum_{h_i \in H} P(v_j|h_i)P(h_i|D)$$

Remark: $h_{MAP}(x)$ is not the most probable classification of x !
(See the next example.)

The Bayes Optimal Classifier: An Example

Let us consider three possible hypotheses:

$$P(h_1|D) = 0.4, \quad P(h_2|D) = 0.3, \quad P(h_3|D) = 0.3$$

Obviously, $h_{MAP} = h_1$.

Let's consider an instance x such that

$$h_1(x) = +, \quad h_2(x) = -, \quad h_3(x) = -$$

Question: What is the most probable classification of x ?

Answer:

$$P(-|h_1) = 0, \quad P(+|h_1) = 1$$

$$P(-|h_2) = 1, \quad P(+|h_2) = 0$$

$$P(-|h_3) = 1, \quad P(+|h_3) = 0$$

$$\sum_{h_i \in H} P(+|h_i)P(h_i|D) = 0.4 \quad \text{and} \quad \sum_{h_i \in H} P(-|h_i)P(h_i|D) = 0.6$$

therefore

$$\operatorname{argmax}_{v_j \in V} \sum_{h_i \in H} P(v_j|h_i)P(h_i|D) = -$$

2.3 The Gibbs Classifier

[Oppor and Haussler, 1991]

Note: The Bayes optimal classifier provides the best result, but it can be expensive if there are many hypotheses.

Gibbs algorithm:

1. Choose one hypothesis at random, according to $P(h|D)$
2. Use this to classify new instance

Surprising fact [Haussler et al. 1994]:

If the target concept is selected randomly according to the $P(h|D)$ distribution, then the expected error of Gibbs Classifier is no worse than twice the expected error of the Bayes optimal classifier!

$$E[\text{error}_{\text{Gibbs}}] \leq 2E[\text{error}_{\text{BayesOptimal}}]$$

2.4 The Naive Bayes Classifier

When to use it:

- The target function f takes value from a finite set $V = \{v_1, \dots, v_k\}$
- Moderate or large training data set is available
- The attributes $\langle a_1, \dots, a_n \rangle$ that describe instances are conditionally independent w.r.t. to the given classification:

$$P(a_1, a_2 \dots a_n | v_j) = \prod_i P(a_i | v_j)$$

The most probable value of $f(x)$ is:

$$\begin{aligned} v_{MAP} &= \operatorname{argmax}_{v_j \in V} P(v_j | a_1, a_2 \dots a_n) = \operatorname{argmax}_{v_j \in V} \frac{P(a_1, a_2 \dots a_n | v_j) P(v_j)}{P(a_1, a_2 \dots a_n)} \\ &= \operatorname{argmax}_{v_j \in V} P(a_1, a_2 \dots a_n | v_j) P(v_j) = \operatorname{argmax}_{v_j \in V} \prod_i P(a_i | v_j) P(v_j) \stackrel{\text{not.}}{=} v_{NB} \end{aligned}$$

This is the so-called *decision rule* of the Naive Bayes classifier.

The Joint Bayes Classifier

$$\begin{aligned} v_{MAP} &= \operatorname{argmax}_{v_j \in V} P(v_j | a_1, a_2 \dots a_n) = \dots \\ &= \operatorname{argmax}_{v_j \in V} P(a_1, a_2 \dots a_n | v_j) P(v_j) = \operatorname{argmax}_{v_j \in V} P(a_1, a_2 \dots a_n, v_j) \stackrel{not.}{=} v_{JB} \end{aligned}$$

The Naive Bayes Classifier: Remarks

1. Along with decision trees, neural networks, k-nearest neighbours, the Naive Bayes Classifier is **one of the most practical** learning methods.
2. Compared to the previously presented learning algorithms, the Naive Bayes Classifier **does no search** through the hypothesis space;
the output hypothesis is simply formed by estimating the **parameters** $P(v_j)$, $P(a_i|v_j)$.

The Naive Bayes Classification Algorithm

NAIVE_BAYES_LEARN(*examples*)

 for each target value v_j

$\hat{P}(v_j) \leftarrow \text{estimate } P(v_j)$

 for each attribute value a_i of each attribute a

$\hat{P}(a_i|v_j) \leftarrow \text{estimate } P(a_i|v_j)$

CLASSIFY_NEW_INSTANCE(x)

$v_{NB} = \operatorname{argmax}_{v_j \in V} \hat{P}(v_j) \prod_{a_i \in x} \hat{P}(a_i|v_j)$

The Naive Bayes: An Example

Consider again the *PlayTennis* example, and new instance

$\langle Outlook = sun, Temp = cool, Humidity = high, Wind = strong \rangle$

We compute:

$$v_{NB} = \operatorname{argmax}_{v_j \in V} P(v_j) \prod_i P(a_i | v_j)$$

$$P(yes) = \frac{9}{14} = 0.64 \quad P(no) = \frac{5}{14} = 0.36$$

...

$$P(strong|yes) = \frac{3}{9} = 0.33 \quad P(strong|no) = \frac{3}{5} = 0.60$$

$$P(yes) P(sun|yes) P(cool|yes) P(high|yes) P(strong|yes) = 0.0053$$

$$P(no) P(sun|no) P(cool|no) P(high|no) P(strong|no) = 0.0206$$

$$\rightarrow v_{NB} = no$$

A Note on The Conditional Independence Assumption of Attributes

$$P(a_1, a_2 \dots a_n | v_j) = \prod_i P(a_i | v_j)$$

It is often violated in practice ...but it works surprisingly well anyway.

Note that we don't need estimated posteriors $\hat{P}(v_j|x)$ to be correct; we only need that

$$\operatorname{argmax}_{v_j \in V} \hat{P}(v_j) \prod_i \hat{P}(a_i | v_j) = \operatorname{argmax}_{v_j \in V} P(v_j) P(a_1 \dots, a_n | v_j)$$

[Domingos & Pazzani, 1996] analyses this phenomenon.

Naive Bayes Classification: The problem of unseen data

What if none of the training instances with target value v_j have the attribute value a_i ?

It follows that $\hat{P}(a_i|v_j) = 0$, and $\hat{P}(v_j) \prod_i \hat{P}(a_i|v_j) = 0$

The typical **solution** is to (re)define $P(a_i|v_j)$, for each value v_j of a_i :

$\hat{P}(a_i|v_j) \leftarrow \frac{n_c + mp}{n + m}$, where

- n is number of training examples for which $v = v_j$,
- n_c number of examples for which $v = v_j$ and $a = a_i$
- p is a prior estimate for $\hat{P}(a_i|v_j)$
(for instance, if the attribute a has k values, then $p = \frac{1}{k}$)
- m is a weight given to that prior estimate
(i.e. number of “virtual” examples)

Using the Naive Bayes Learner: Learning to Classify Text

- Learn which news articles are of interest

Target concept *Interesting?* : *Document* $\rightarrow \{+, -\}$

- Learn to classify web pages by topic

Target concept *Category* : *Document* $\rightarrow \{c_1, \dots, c_n\}$

Naive Bayes is among most effective algorithms

Learning to Classify Text: Main Design Issues

1. Represent each document by a vector of words

- one attribute per word position in document

2. Learning:

- use training examples to estimate $P(+)$, $P(-)$, $P(doc|+)$, $P(doc|-)$
- Naive Bayes conditional independence assumption:

$$P(doc|v_j) = \prod_{i=1}^{length(doc)} P(a_i = w_k|v_j)$$

where $P(a_i = w_k|v_j)$ is probability that word in position i is w_k , given v_j

- Make one more assumption:

$$\forall i, m \ P(a_i = w_k|v_j) = P(a_m = w_k|v_j) = P(w_k|v_j)$$

i.e. attributes are (not only indep. but) also identically distributed

LEARN_NAIVE_BAYES_TEXT(*Examples*, *Vocabulary*)

1. Collect all words and other tokens that occur in *Examples*

Vocabulary \leftarrow all distinct words and other tokens in *Examples*

2. Calculate the required $P(v_j)$ and $P(w_k|v_j)$ probability terms

For each target value v_j in V

$docs_j \leftarrow$ the subset of *Examples* for which the target value is v_j

$$P(v_j) \leftarrow \frac{|docs_j|}{|Examples|}$$

$Text_j \leftarrow$ a single doc. created by concat. all members of $docs_j$

$n \leftarrow$ the total number of words in $Text_j$

For each word w_k in *Vocabulary*

$n_k \leftarrow$ the number of times word w_k occurs in $Text_j$

$$P(w_k|v_j) \leftarrow \frac{n_k+1}{n+|Vocabulary|} \quad (\text{here we use the } m\text{-estimate})$$

CLASSIFY_NAIVE_BAYES_TEXT(*Doc*)

positions \leftarrow all word positions in *Doc* that contain tokens from *Vocabulary*

Return $v_{NB} = \operatorname{argmax}_{v_j \in V} P(v_j) \prod_{i \in \text{positions}} P(a_i = w_k | v_j)$


Application: Learning to Classify Usenet News Articles

Given 1000 training documents from each of the 20 newsgroups, learn to classify new documents according to which newsgroup it came from

| | |
|--------------------------|--------------------|
| comp.graphics | misc.forsale |
| comp.os.ms-windows.misc | rec.autos |
| comp.sys.ibm.pc.hardware | rec.motorcycles |
| comp.sys.mac.hardware | rec.sport.baseball |
| comp.windows.x | rec.sport.hockey |
| alt.atheism | sci.space |
| soc.religion.christian | sci.crypt |
| talk.religion.misc | sci.electronics |
| talk.politics.mideast | sci.med |
| talk.politics.misc | |
| talk.politics.guns | |

Naive Bayes: 89% classification **accuracy** (having used 2/3 of each group for training; eliminated rare words, and the 100 most freq. words)

Learning Curve for 20 Newsgroups



2.5 The Minimum Description Length Principle

Occam's razor: prefer the shortest hypothesis

Bayes analysis: prefer the hypothesis h_{MAP}

$$\begin{aligned} h_{MAP} &= \operatorname{argmax}_{h \in H} P(D|h)P(h) = \operatorname{argmax}_{h \in H} (\log_2 P(D|h) + \log_2 P(h)) \\ &= \operatorname{argmin}_{h \in H} (-\log_2 P(D|h) - \log_2 P(h)) \end{aligned}$$

Interesting fact from the Information Theory:

The **optimal** (shortest expected coding length) **code** for an event with probability p is the one using $-\log_2 p$ bits.

So we can interpret:

$-\log_2 P(h)$: the length of h under the optimal code

$-\log_2 P(D|h)$: the length of D given h under the optimal code

Therefore we prefer the hypothesis h that minimizes...

Bayes Analysis and the MDL Principle

We saw that a MAP learner prefers the hypothesis h that minimizes $L_{C_1}(h) + L_{C_2}(D|h)$, where $L_C(x)$ is the description length of x under encoding C

$$h_{MDL} = \operatorname{argmin}_{h \in H} (L_{C_1}(h) + L_{C_2}(D|h))$$

Example: H = decision trees, D = training data labels

- $L_{C_1}(h)$ is the number of bits to describe tree h
- $L_{C_2}(D|h)$ is the number of bits to describe D given h

In literature, the application of MDL to practical problems often include arguments justifying the choice of the encodings C_1 and C_2 .

For instance:

$L_{C_2}(D|h) = 0$ if examples are classified perfectly by h ,
and both the transmitter and the receiver know h .

Therefore, in this situation we need only to describe exceptions. So:

$$h_{MDL} = \operatorname{argmin}_{h \in H} (\operatorname{length}(h) + \operatorname{length}(\operatorname{misclassifications}))$$

In general, MDL trades off hypothesis size for training errors:

it might select a shorter hypothesis that makes few errors over a longer hypothesis that perfectly classifies the data!

Consequence: In learning (for instance) decision trees, (using) the MDL principle can work as an alternative to pruning.

The MDL Principle: Back to Occam's Razor


MDL hypotheses are not necessarily also the best/MAP ones.

(For that, we should know all the probabilities $P(D|h)$ and $P(h)$.)

3 Learning Maximum Likelihood (ML) Hypothesis

3.1 Learning Real Valued Functions:

ML Hypotheses as Least Squared Error Hypotheses



Problem: Consider learning a real-valued target function $f : X \rightarrow \mathbb{R}$ from D , a training set consisting of examples $\langle x_i, d_i \rangle, i = 1, \dots, m$ with

x_i , assumed fixed (to simplify)

d_i noisy training value $d_i = f(x_i) + e_i$

e_i is random variable (noise) drawn independently for each x_i , according to some Gaussian distribution with mean=0.

Proposition

Considering H , a certain class of functions $h : X \rightarrow \mathbb{R}$ such that $h(x_i) = f(x_i)$ and assuming that x_i are mutually independent given h ,
the maximum likelihood hypothesis h_{ML} is the one that minimizes the sum of squared errors:

$$h_{ML} \stackrel{\text{def.}}{=} \operatorname{argmax}_{h \in H} P(D|h) = \operatorname{argmin}_{h \in H} \sum_{i=1}^m (d_i - h(x_i))^2$$

Proof

30.

Note: We will use the **probability density function**:

$$p(x_0) \stackrel{def.}{=} \lim_{\epsilon \rightarrow 0} \frac{1}{\epsilon} P(x_0 \leq x < x_0 + \epsilon)$$

$$\begin{aligned} h_{ML} &= \operatorname{argmax}_{h \in H} P(D|h) = \operatorname{argmax}_{h \in H} \prod_{i=1}^m p(d_i|h) \stackrel{\mu_i = f(x_i)}{=} \operatorname{argmax}_{h \in H} \prod_{i=1}^m p(e_i|h) \\ &= \operatorname{argmax}_{h \in H} \prod_{i=1}^m p(d_i - f(x_i)|h) \stackrel{h(x_i) = f(x_i)}{=} \operatorname{argmax}_{h \in H} \prod_{i=1}^m p(d_i - h(x_i)|h) \\ &= \operatorname{argmax}_{h \in H} \prod_{i=1}^m \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2}\left(\frac{d_i - h(x_i)}{\sigma}\right)^2} = \operatorname{argmax}_{h \in H} \left(\sum_{i=1}^m \ln \frac{1}{\sqrt{2\pi\sigma^2}} - \frac{1}{2} \left(\frac{d_i - h(x_i)}{\sigma} \right)^2 \right) \\ &= \operatorname{argmax}_{h \in H} \sum_{i=1}^m -\frac{1}{2} \left(\frac{d_i - h(x_i)}{\sigma} \right)^2 = \operatorname{argmax}_{h \in H} \sum_{i=1}^m -(d_i - h(x_i))^2 \\ &= \operatorname{argmin}_{h \in H} \sum_{i=1}^m (d_i - h(x_i))^2 \end{aligned}$$

Generalisations...

1. Similar derivations can be performed starting with **other assumed noise distributions** (than Gaussians), producing **different results**.
2. It was assumed that
 - a.* the **noise** affects only $f(x_i)$, and
 - b.* no noise was recorded in the **attribute values** for the given examples x_i .

Otherwise, the analysis becomes significantly more complex.

3.2 ML hypotheses for Learning Probability Functions

Let us consider a non-deterministic function (i.e. one-to-many relation) $f : X \rightarrow \{0, 1\}$.

Given a set of independently drawn examples

$D = \{ \langle x_1, d_1 \rangle, \dots, \langle x_m, d_m \rangle \}$ where $d_i = f(x_i) \in \{0, 1\}$,

we would like to learn a ML hypothesis for the probability function $g(x) \stackrel{\text{def.}}{=} P(f(x) = 1)$.

For example, $h(x_i) = 0.92$ if $P(\{ \langle x_i, d_i \rangle \mid d_i = 1 \}) = 0.92$.

Proposition: In this setting, $h_{ML} = \underset{h \in H}{\operatorname{argmax}} P(D \mid h)$ maximizes the sum $\sum_{i=1}^m [d_i \ln h(x_i) + (1 - d_i) \ln (1 - h(x_i))]$.

Proof:

$$P(D \mid h) = \prod_{i=1}^m P(x_i, d_i \mid h) = \prod_{i=1}^m P(d_i \mid x_i, h) \cdot P(x_i \mid h)$$

It can be assumed that x_i is independent of h , therefore:

$$P(D \mid h) = \prod_{i=1}^m P(d_i \mid x_i, h) \cdot P(x_i)$$

Proof (continued):

What we wanted to compute is $h(x_i) = P(d_i = 1 \mid x_i, h)$.

In a more general form:

$$P(d_i \mid x_i, h) = \begin{cases} h(x_i) & \text{if } d_i = 1 \\ 1 - h(x_i) & \text{if } d_i = 0 \end{cases}$$

In a more convenient mathematical form: $P(d_i \mid x_i, h) = h(x_i)^{d_i} (1 - h(x_i))^{1-d_i}$.

$$\begin{aligned} \Rightarrow h_{ML} &= \mathbf{argmax}_{h \in H} \prod_{i=1}^m [h(x_i)^{d_i} (1 - h(x_i))^{1-d_i} P(x_i)] \\ &= \mathbf{argmax}_{h \in H} \prod_{i=1}^m h(x_i)^{d_i} (1 - h(x_i))^{1-d_i} \cdot \prod_{i=1}^m P(x_i) \\ &= \mathbf{argmax}_{h \in H} \prod_{i=1}^m h(x_i)^{d_i} (1 - h(x_i))^{1-d_i} \\ &= \mathbf{argmax}_{h \in H} \sum_{i=1}^m [d_i \ln h(x_i) + (1 - d_i) \ln (1 - h(x_i))] \end{aligned}$$

Note: The quantity $-\sum_{i=1}^m [d_i \ln h(x_i) + (1 - d_i) \ln (1 - h(x_i))]$ is called cross-entropy; the above h_{ML} minimizes this quantity.

3.3 The Expectation Maximization (EM) Algorithm

[Dempster et al, 1977]

Find (local) Maximum Likelihood hypotheses when
data is only partially observable:

- Unsupervised learning (i.e., clustering):
the target value is unobservable
- Supervised learning:
some instance attributes are unobservable

Some applications:

- Non-hierarchical clustering:
Estimate the means of k Gaussians
- Learn Hidden Markov Models
- Learn Probabilistic Context Free Grammars
- Train Radial Basis Function Networks
- Train Bayesian Belief Networks

The General EM Problem

Given

- observed data $X = \{x_1, \dots, x_m\}$
independently generated using the parameterized distributions/hypotheses h_1, \dots, h_m
- unobserved data $Z = \{z_1, \dots, z_m\}$

determine

\hat{h} that (locally) maximizes $P(Y|h)$,
where $Y = \{y_1, \dots, y_m\}$ is the full data $y_i = x_i \cup z_i$

The Essence of the EM Approach

Start with h^0 , an arbitrarily/conveniently chosen value of h .

Repeatedly

1. Use the observed data X and the current hypothesis h^t to **estimate [the probabilities associated to the values of] the unobserved** variables Z , and further on compute their expectations, $E[Z]$.
2. The expected values of the unobserved variables Z are used to **calculate an improved hypothesis** h^{t+1} , based on maximizing the mean of a log-verosimilarity function: $E[\ln P(Y|h)|X, h^t]$.

The General EM Algorithm

Repeat the following two steps until convergence is reached:

Estimation (E) step:

Calculate the log **likelihood function**


$$Q(h|h^t) \stackrel{not.}{=} E[\ln P(Y|h)|X, h^t]$$

where $Y = X \cup Z$.

Maximization (M) step:

Replace hypothesis h^t by the hypothesis h^{t+1} that maximizes this Q function.

$$h^{t+1} \leftarrow \operatorname{argmax}_h Q(h|h^t)$$



Baum-Welch Theorem

When Q is continuous, it can be shown that EM converges to a stationary point (local maximum) of the likelihood function $P(Y|h)$.

4 Bayesian Belief Networks

(also called Bayes Nets)

Interesting because:

- The Naive Bayes assumption of conditional independence of attributes is too restrictive.
(But it's intractable without some such assumptions...)
- Bayesian Belief networks describe **conditional independence among *subsets* of variables**.
- It allows the combination of prior knowledge about (in)dependencies among variables with observed training data.

Conditional Independence

Definition: X is **conditionally independent** of Y given Z if the probability distribution governing X is independent of the value of Y given a value of Z :

$$(\forall x_i, y_j, z_k) \ P(X = x_i | Y = y_j, Z = z_k) = P(X = x_i | Z = z_k)$$

More compactly, we write $P(X|Y, Z) = P(X|Z)$

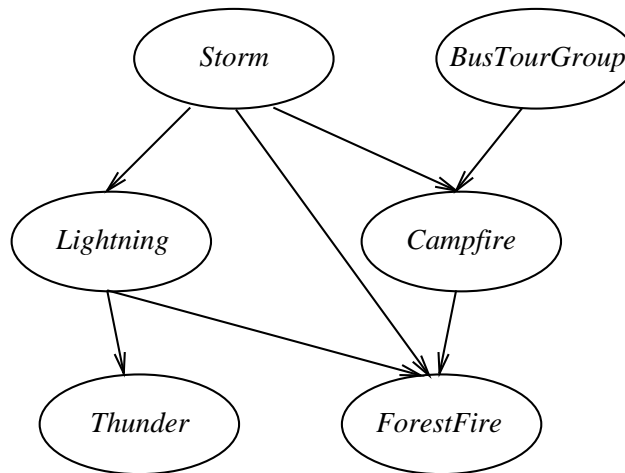
Note: Naive Bayes uses conditional independence to justify

$$P(A_1, A_2 | V) = P(A_1 | A_2, V) P(A_2 | V) = P(A_1 | V) P(A_2 | V)$$

Generalizing the above definition:

$$P(X_1 \dots X_l | Y_1 \dots Y_m, Z_1 \dots Z_n) = P(X_1 \dots X_l | Z_1 \dots Z_n)$$

A Bayes Net



| | S, B | $S, \neg B$ | $\neg S, B$ | $\neg S, \neg B$ |
|----------|--------|-------------|-------------|------------------|
| C | 0.4 | 0.1 | 0.8 | 0.2 |
| $\neg C$ | 0.6 | 0.9 | 0.2 | 0.8 |



The network is defined by

- A directed acyclic graph, representing a set of conditional independence assertions:

Each node — representing a random variable — is asserted to be conditionally independent of its nondescendants, given its immediate predecessors.

Example: $P(\text{Thunder} | \text{ForestFire}, \text{Lightning}) = P(\text{Thunder} | \text{Lightning})$

- A table of local conditional probabilities for each node/variable.

A Bayes Net (Cont'd)

represents **the joint probability distribution** over all variables Y_1, Y_2, \dots, Y_n :

This joint distribution is fully defined by the graph, plus the conditional probabilities:

$$P(y_1, \dots, y_n) = P(Y_1 = y_1, \dots, Y_n = y_n) = \prod_{i=1}^n P(y_i | Parents(Y_i))$$

where $Parents(Y_i)$ denotes immediate predecessors of Y_i in the graph.

In our **example**: $P(Storm, BusTourGroup, \dots, ForestFire)$

Inference in Bayesian Nets


Question: Given a Bayes net, can one infer the probabilities of values of one or more network variables, given the observed values of (some) others?

Example:

Given the Bayes net

compute:

- (a) $P(S)$
- (b) $P(A, S)$
- (b) $P(A)$



Inference in Bayesian Nets (Cont'd)

Answer(s):

- If only one variable is of unknown (probability) value, then it is easy to infer it
- In the general case, we can compute the probability distribution for any subset of network variables, given the distribution for any subset of the remaining variables. But...
- The exact inference of probabilities for an arbitrary Bayes net is an NP-hard problem!!

Inference in Bayesian Nets (Cont'd)

In practice, we can succeed in many cases:

- Exact inference methods work well for some net structures.
- Monte Carlo methods “simulate” the network randomly to calculate approximate solutions [Pradham & Dagum, 1996].

(In theory even approximate inference of probabilities in Bayes Nets can be NP-hard!! [Dagum & Luby, 1993])

Learning Bayes Nets (I)

There are **several variants** of this learning task

- The network structure might be either *known* or *unknown* (i.e., it has to be inferred from the training data).
- The training examples might provide values of *all* network variables, or just for *some* of them.

The simplest case:

If the structure is known and we can observe the values of all variables,
then it is easy to estimate the conditional probability table entries (analogous to training a Naive Bayes classifier).

Learning Bayes Nets (II)

When

- the structure of the Bayes Net is known, and
- the variables are only partially observable in the training data

learning the entries in the conditional probabilities tables is similar to (learning the weights of hidden units in) training a neural network with hidden units:

- We can learn the net's conditional probability tables using the gradient ascent!
- Converge to the network h that (locally) maximizes $P(D|h)$.

Gradient Ascent for Bayes Nets

49.

Let w_{ijk} denote one entry in the conditional probability table for the variable Y_i in the network

$$w_{ijk} = P(Y_i = y_{ij} | Parents(Y_i) = \text{the list } u_{ik} \text{ of values})$$

It can be shown (see the next two slides) that

$$\frac{\partial \ln P_h(D)}{\partial w_{ijk}} = \sum_{d \in D} \frac{P_h(y_{ij}, u_{ik} | d)}{w_{ijk}}$$

therefore perform gradient ascent by repeatedly

1. **update all** w_{ijk} using the training data D

2. **renormalize the** w_{ijk} to assure

$$w_{ijk} \leftarrow w_{ijk} + \eta \sum_{d \in D} \frac{P_h(y_{ij}, u_{ik} | d)}{w_{ijk}}$$

$$\sum_j w_{ijk} = 1 \text{ and } 0 \leq w_{ijk} \leq 1$$

Gradient Ascent for Bayes Nets: Calculus

$$\frac{\partial \ln P_h(D)}{\partial w_{ijk}} = \frac{\partial}{\partial w_{ijk}} \ln \prod_{d \in D} P_h(d) = \sum_{d \in D} \frac{\partial \ln P_h(d)}{\partial w_{ijk}} = \sum_{d \in D} \frac{1}{P_h(d)} \frac{\partial P_h(d)}{\partial w_{ijk}}$$

Summing over all values $y_{ij'}$ of Y_i , and $u_{ik'}$ of $U_i = \text{Parents}(Y_i)$:

$$\begin{aligned} \frac{\partial \ln P_h(D)}{\partial w_{ijk}} &= \sum_{d \in D} \frac{1}{P_h(d)} \frac{\partial}{\partial w_{ijk}} \sum_{j'k'} P_h(d|y_{ij'}, u_{ik'}) P_h(y_{ij'}, u_{ik'}) \\ &= \sum_{d \in D} \frac{1}{P_h(d)} \frac{\partial}{\partial w_{ijk}} \sum_{j'k'} P_h(d|y_{ij'}, u_{ik'}) P_h(y_{ij'}|u_{ik'}) P_h(u_{ik'}) \end{aligned}$$

Note that $w_{ijk} \equiv P_h(y_{ij}|u_{ik})$, therefore...

Gradient Ascent for Bayes Nets: Calculus (Cont'd)

$$\begin{aligned}
 \frac{\partial \ln P_h(D)}{\partial w_{ijk}} &= \sum_{d \in D} \frac{1}{P_h(d)} \frac{\partial}{\partial w_{ijk}} P_h(d|y_{ij}, u_{ik}) w_{ijk} P_h(u_{ik}) \\
 &= \sum_{d \in D} \frac{1}{P_h(d)} P_h(d|y_{ij}, u_{ik}) P_h(u_{ik}) \quad (\text{applying Bayes th.}) \\
 &= \sum_{d \in D} \frac{1}{P_h(d)} \frac{P_h(y_{ij}, u_{ik}|d) P_h(d) P_h(u_{ik})}{P_h(y_{ij}, u_{ik})} \\
 &= \sum_{d \in D} \frac{P_h(y_{ij}, u_{ik}|d) P_h(u_{ik})}{P_h(y_{ij}, u_{ik})} = \sum_{d \in D} \frac{P_h(y_{ij}, u_{ik}|d)}{P_h(y_{ij}|u_{ik})} \\
 &= \sum_{d \in D} \frac{P_h(y_{ij}, u_{ik}|d)}{w_{ijk}}
 \end{aligned}$$

Learning Bayes Nets (II, Cont'd)

The **EM** algorithm can also be used.

Repeatedly:

1. Calculate/estimate from data the probabilities of unobserved variables w_{ijk} ,
assuming that the hypothesis h holds
2. Calculate a new h (i.e. new values of w_{ijk}) so to maximize $E[\ln P(D|h)]$,
where D now includes both the observed and the unobserved variables.

Learning Bayes Nets (III)

When the **structure is unknown**, algorithms usually use greedy search to trade off network complexity (add/subtract edges/nodes) against degree of fit to the data.

Example: [Cooper & Herscovitz, 1992] the *K2 algorithm*:

When data is fully observable, use a score metric to choose among alternative networks.

They report an experiment on (re-learning) a network with 37 nodes and 46 arcs describing anesthesia problems in a hospital operating room. Using 3000 examples, the program succeeds almost perfectly: it misses one arc and adds an arc which is not in the original net.