

# Reinforcement Learning

## Framework

Mario Martin  
Universitat politècnica de Catalunya  
Dept. LSI

# Motivation

- Target is to learn a behavior instead of learning a concept
- Agent discovers the behavior
  - No examples required for learning
  - No need of an expert to know beforehand the behavior to be learnt
- Learning is grounded in the environment of the agent
- No bias in learning.
- Agent learns the *optimal* behavior.

# Motivation

Reinforcement Learning captures  
essential features of agent learning

# Costs

- Some assumptions are required
- Long computational time and a lot of memory required
- Fortunately these inconveniences can be reduced in some cases.


# Definition of Reinforcement Learning

1. Learning about, from, and while interacting with an environment to achieve a goal

*... read as ...*

2. Learning a mapping from situations to actions to maximize long-term reward, without using a model of the world

# The Agent-Environment Interaction



Agent and environment interact at discrete time steps:  $t = 0, 1, 2, \dots$


Agent observes state at step  $t$ :  $s_t \in S$

produces action at step  $t$ :  $a_t \in A(s_t)$

gets resulting reward:  $r_{t+1} \in \mathfrak{R}$

and resulting next state:  $s_{t+1}$

# Trace of a trial



# Markovian Decision Process (MDP)

- Agent-env. interaction becomes a MDP when
  - Finite set of situations
  - Finite set of actions
  - Transition probabilities:

$$P_{ss'}^a = \Pr \{s_{t+1} = s' \mid s_t = s, a_t = a\} \quad \text{for all } s, s' \in S, a \in A(s).$$

- Reward probabilities:


$$R_{ss'}^a = E \{r_{t+1} \mid s_t = s, a_t = a, s_{t+1} = s'\} \quad \text{for all } s, s' \in S, a \in A(s).$$



# The Markovian Property

- When actions and states are not finite, discretize the set of actions and states
- When transition probabilities do not depend only on the current state,
  - The agent could represent states as structures built up over time from sequences of sensations (for instance, for the agent the state is a window of the recent states)
  - Use POMDP learning algorithms

# Elements of RL



- **Policy:** what to do
- **Reward:** what is good
- **Value:** what is good because it *predicts* reward
- **Model:** what follows what

# Components of an RL Agent (I)

- Policy (behavior)
  - Mapping from states to actions

$$\pi(s) \longrightarrow a$$

- Reward
  - Local reward in state  $t$

$$r_t$$

# Components of an RL Agent (II)

- Model
  - $T(s,a,s')$  : probability of transition from state  $s$  to  $s'$  executing action  $a$
- The transition probabilities depends only on these parameters (Markovian model)
- Not known by the agent

# Components of an RL Agent (III)


- Value functions
  - $V^\pi(s)$  : Long-term reward estimation from state  $s$  following policy  $\pi$
  - $Q^\pi(s, a)$  : Long-term reward estimation from state  $s$  executing action  $a$  and then following policy  $\pi$

# Simple Example

- Markovian model, reward function, set of actions, set of states
- Maze:

|   |    |  |  |      |     |
|---|----|--|--|------|-----|
|   |    |  |  | -100 | +10 |
|   |    |  |  |      |     |
| A | -1 |  |  |      |     |
|   |    |  |  |      |     |
|   |    |  |  |      |     |
|   |    |  |  |      |     |

Actions, 90% reliable



# Getting the Degree of Abstraction Right

- Actions can be low level (e.g., voltages to motors), or high level (e.g., accept a job offer), “mental” (e.g., shift in focus of attention), etc.
- States can low-level “sensations”, or they can be abstract, symbolic, based on memory, or subjective (e.g., the state of being “surprised” or “lost”).
- An RL agent is not like a whole animal or robot, which consist of many RL agents as well as other components.
- Reward computation is in the agent’s environment because the agent cannot change it arbitrarily.